

1. Set configuration context `$ kubectl config use-context k8s`

Monitor the logs of Pod `foobar` and

- Extract log lines corresponding to error `unable-to-access-website`
- Write them to `/opt/KULM00201/foobar`

Question weight 5%

Answer:

```
$ kubectl logs pod foobar | grep "unable-to-access-website" >
/opt/KULM00201/foobar
```

2. Set configuration context `$ kubectl config use-context k8s`

List all PVs sorted by **name**, saving the full `kubectl` output to `/opt/KUCC0010/my_volumes`. Use `kubectl`'s own functionality for sorting the output, and do not manipulate it any further.

Question weight 3%

Answer:

```
$ kubectl get pv --all-namespaces --sort-by="metadata.name" >>
/opt/KUCC0010/my_volumes or
$ kubectl get pv --all-namespaces --sort-by="spec.capacity.storage" >>
/opt/KUCC0010/my_volumes
```

3. Set configuration context `$ kubectl config use-context k8s`

Ensure a single instance of Pod `nginx` is running on each node of the kubernetes cluster where `nginx` also represents the image name which has to be used. Do not override any taints currently in place.

Use **Daemon sets** to complete this task and use `ds.kusc00201` as Daemonset name.

Question weight 3%

Answer:

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: ds.kusc...
spec:
  selector:
```

```

matchLabels:
  name: nginx
template:
  metadata:
    labels:
      name: nginx
  spec:
    containers:
      - name: nginx
        image: nginx

```

4. Set configuration context \$ **kubectrl config use-context k8s**

Perform the following tasks

- Add an init container to **lumpy--koala** (Which has been defined in spec file **/opt/kucc00100/pod-spec-KUCC00100.yaml**)
- The init container should create an empty file named **/workdir/calm.txt**
- If **/workdir/calm.txt** is not detected, the Pod should exit
- Once the spec file has been updated with the init container definition, the Pod should be created.

Question weight 7%

yaml 文件:

apiVersion: v1

kind: Pod

Metadate:

 Name lumpy-koala

Spec

 Volumes

 - name workdir

 Emptydir

 Containers

 - name checker

 Image alpine

 Command

 Volumemount

 - name workdir

 Mountpath /workdir

Answer:

apiVersion: v1

kind: pod

metadata:

 name: lumpy-koala

```
spec:
  volumes:
    - name: workdir
      emptyDir: {}
  containers:
    - name: checker
      image: alpine
      command: ....
      volumeMount:
        - name: workdir
          mountPath: /workdir
```

initContainers:

```
- name: init-lumpy-koala
  image: busybox
  command: ['sh','-c','touch /workdir/calm.txt']
  volumeMounts:
    - name: workdir
      mountPath: "/workdir"
```

kubectl create -f `opt/kucc00100/pod-spec-KUCC00100.yaml`

5. Set configuration context `$ kubectl config use-context k8s`

Create a pod named `kucc4` with a single container for each of the following images running inside (there may be between 1 and 4 images specified): `nginx + redis + memcached + consul`

Question weight: 4%

Answer:

apiVersion: v1

kind: Pod

metadata:

name: `kucc4`

spec:

containers:

```
- name: nginx
  image: nginx
```

containers:

```
- name: redis
  image: redis
```

containers:

```
- name: memcached
  image: memcached
```

containers:

- **name: consul**
- image: consul**

kubectl create -f kucc4.yaml

6. Set configuration context **\$ kubectl config use-context k8s**

Schedule a Pod as follows:

- Name: **nginx-kusc00101**
- Image: **nginx**
- Node selector: **disk=ssd**

Question weight: 2%

Answer:

apiVersion: v1

kind: Pod

metadata:

name: nginx-kusc00101

spec:

containers:

- **name: nginx-kusc00101**
- image: nginx**

nodeSelector:

disk: ssd

kubectl create -f nginx-kusc00101.yaml

7. Set configuration context **\$ kubectl config use-context k8s**

Create a deployment as follows

- Name: **nginx-app**
- Using container **nginx** with version **1.11.9-alpine**
- The deployment should contain **3** replicas

Next, deploy the app with new version **1.12.0-alpine** by performing a rolling update and record that update.

Finally, rollback that update to the previous version **1.11.9-alpine**

Question weight: 4%

Answer:

kubectl run nginx-app --image='nginx:1.11.9-alpine'
--replicas=3

```
kubectl set image deployment/nginx-app  
nginx-app='nginx:1.12.0-alpine' --record=true
```

```
kubectl rollout undo deployment/nginx-app
```

8. Set configuration context `$ kubectl config use-context k8s`

Create and configure the service `front-end-service` so it's accessible through `NodePort` / `ClusterIP` and routes to the existing pod named `front-end`

Question weight: 4%

Answer:

```
kubectl expose pod front-end --name=front-end-service  
--type="NodePort" --port=80
```

Or

9. Set configuration context `$ kubectl config use-context k8s`

Create a Pod as follows:

- Name: `jenkins`
- Using image: `jenkins`
- In a new Kubernetes namespace named `website-frontend`

Question weight 3%

Answer:

```
kubectl create ns website-frontend
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: jenkins
```

```
  namespace: website-frontend
```

```
spec:
```

```
  containers:
```

```
    - name: jenkins
```

```
      image: jenkins
```

```
kubectl create -f jenkins.yaml
```

10. Set configuration context `$ kubectl config use-context k8s`

Create a deployment spec file that will:

- Launch 7 replicas of the **redis** image with the label: **app_env_stage=dev**
- Deployment name: **kual00201**

Save a copy of this spec file to **/opt/KUAL00201/deploy_spec.yaml** (or **.json**)

When you are done, clean up (delete) any new k8s API objects that you produced during this task

Question weight: 3%

Answer:

apiVersion: apps/v1

kind: Deployment

metadata:

name: kual00201

spec:

replicas: 7

selector:

matchLabels:

app_env_stage: dev

template:

metadata:

labels:

app_env_stage: dev

spec:

containers:

- name: redis

image: redis

kubectl create -f kual00201.yaml

cp kual00201.yaml /opt/KUAL00201/deploy_spec.yaml

kubectl delete deploy kual00201

11. Set configuration context **\$ kubectl config use-context k8s**

Create a file **/opt/KUCC00302/kucc00302.txt** that lists all pods that implement Service **foo** in Namespace **production**.

The format of the file should be one pod name per line.

Question weight: 3%

Answer:

kubectl get svc -o wide -n production

得到 label

kubectl get pod -l label=value | awk '{print \$1}' | grep -v "NAME" >> /opt/KUCC00302/kucc...

kubectl get po -l label

12. Set configuration context `$ kubectl config use-context k8s`

Create a Kubernetes Secret as follows:

- Name: `super-secret`
- Credential: `alice` or username: `bob`

Create a Pod named `pod-secrets-via-file` using the `redis` image which mounts a secret named `super-secret` at `/secrets`

Create a second Pod named `pod-secrets-via-env` using the `redis` image, which exports `credential` / `username` as `TOPSECRET / CREDENTIALS`

Question weight: 9%

Answer:

```
kubectl create secret generic spuer-secret
--from-literal=Credential=alice
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: pod-secrets-via-file
```

```
spec:
```

```
  containers:
```

```
    - name: pod-secrets-via-file
```

```
      image: redis
```

```
      volumeMounts:
```

```
        - name: super-secret
```

```
          mountPath: "/secret"
```

```
  volumes:
```

```
    - name: super-secret
```

```
      secret:
```

```
        secretName: super-secret
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: pod-secrets-via-env
```

```
spec:
```

```
  containers:
```

```
    - name: pod-secrets-via-env
```

```
      image: redis
```

```
      env:
```

- name: TOPSECRET
valueFrom:
secretKeyRef:
name: super-secret
key: Credential

13. Set configuration context \$ kubectl config use-context k8s

Create a pod as follows:

- Name: non-persistent-redis
- Container image: redis
- Named-volume with name: cache-control
- Mount path: /data/redis

It should launch in the pre-prod namespace and the volume MUST NOT be persistent.

Question weight: 4%

Answer:

apiVersion: v1
kind: Pod
metadata:
name: non-persistent-redis
namespace: pre-prod
spec:
containers:
- name: non-persistent-redis
image: redis
volumeMounts:
- name: cache-control
mountPath: "/data/redis"
volumes:
- name: cache-control
emptyDir: {}

14. Set configuration context \$ kubectl config use-context k8s

Scale the deployment webserver to 6 pods

Question weight: 1%

Answer:

kubectl scale deployment/webserver --replicas=6

15. Set configuration context \$ kubectl config use-context k8s

Check to see how many nodes are ready (not including nodes tainted NoSchedule) and write the

number to **/opt/.....**

Question weight: 2%

Answer: 2

首先用 **kubectl get node** 查看总共 node 数量，然后 **kubectl describe node** 查看不包含污点值为 NoSchedule 的 node 数量

16. Set configuration context **\$ kubectl config use-context k8s**

From the Pod label **name=cpu-utilizer**, find pods running high CPU workloads and write the name of the Pod consuming most CPU to the file **/opt/.....** (which already exists)

Question weight: 2%

Answer:

通过 **kubectl get pod -l name=cpu-utilizer -o wide** 查看 pod 所在的 node

之后 **kubectl describe node** 查看哪个 pod 请求的 cpu 最多

三个都是 0，我就都写进去了

17. Set configuration context **\$ kubectl config use-context k8s**

Create a deployment as follows

- Name: **nginx-dns**
- Exposed via a service: **nginx-dns**
- Ensure that the service & pod are accessible via their respective DNS records
- The container(s) within any Pod(s) running as a part of this deployment should use the **nginx** image

Next, use the utility **nslookup** to look up the DNS records of the service & pod and write the output to **/opt/.../service.dns** and **/opt/.../pod.dns** respectively. Ensure you use the busybox:1.28 image (or earlier) for any testing, as the latest release has an upstream bug which impacts the use of **nslookup**.

Question weight: 7%

Answer:

kubectl run nginx-dns --image=nginx:latest

kubectl expose deployment/nginx-dns --port=80 --name=nginx-dns

kubectl run busybox -it --rm --image=busybox sh

nslookup nginx-dns #svc 记录

nslookup pod-id-addr #pod 记录

第四行： 是查 service 的 dns

第五行： 是查 pod 的

Nslookup 的结果不能直接输出到文件，我是粘贴到 notepad 上再用

18. No configuration context change required for this item

Create a snapshot of the **etcd** instance running at **https://127.0.0.1:2379** saving the snapshot to the file path **/data/backup/etcd-snapshot.db**

The etcd instance is running **etcd** version 3.2.18

The following TLS certificates/key are supplied for connecting to the server with **etcdctl**

- CA certificate: **/opt/KUCM00302/ca.crt**
- Client certificate: **/opt/KUCM00302/etcd-client.crt**
- Client key: **/opt/KUCM00302/etcd-client.key**

Question weight: 7%

Answer:

```
$ export ETCDCTL_API=3
```

```
$ etcdctl help
```

```
$ etcdctl --endpoints=https://127.0.0.1:2379 \
```

```
--ca-file=/opt/KUCM00302/ca.crt \
```

```
--certfile=/opt/KUCM00302/etcd-client.crt \
```

```
--key=/opt/KUCM00302/etcd-client.key \
```

```
snapshot save /data/backup/etcd-snapshot.db
```

因为 etcd 版本不一样，后面关键字跟的方法也不一样，先 help 看一下，照着帮助写就没问题。

19. Set configuration context **\$ kubectl config use-context ek8s**

Set the node labelled with **name=ek8s-node-1** as unavailable and reschedule all the pods running on it.

Question weight: 4%

Answer:

```
kubectl get node --show-labels |grep name=ek8s-node-1 #找出 node
```

```
kubectl drain node $node
```

20. Set configuration context `$ kubectl config use-context wk8s`

A Kubernetes worker node, labelled with **name=wk8s-node-0** is in state **NotReady** . Investigate why this is the case, and perform any appropriate steps to bring the node to a **Ready** state, ensuring that any changes are made permanent.

Hints:

- You can **ssh** to the failed node using `$ ssh wk8s-node-0`
- You can assume elevated privileges on the node with the following command `$ sudo -i`

Question weight: 4%

Answer:

是 **kubelet** 没有启动

kubectl get node

查看 一个 **node** 是 **notReady**

ssh 上去

systemctl status kubelet

发现没有启动

systemctl start kubelet

21. Set configuration context `$ kubectl config use-context wk8s`

Configure the **kubelet systemd** managed service, on the node labelled with **name=wk8s-node-1**, to launch a Pod containing a single container of image **nginx** named **myservice** automatically. Any spec files required should be placed in the **/etc/kubernetes/manifests** directory on the node.

Hints:

- You can **ssh** to the failed node using `$ ssh wk8s-node-1`
- You can assume elevated privileges on the node with the following command `$ sudo -i`

Question weight: 4%

Answer

1. 登陆到节点上

找到 **kubelet** 的目录

/etc/kubernetes/manifests/

创建下面的 **yaml** , 并保存在这个目录

apiVersion: v1

kind: Pod

metadata:

```
name: webtool
spec:
  containers:
  - name: myservice
    image: nginx
```

2. vi /etc/systemctl/system/kubelet.service

在 /usr/bin/kubelete 下面第一行加入
--pod-manifest-path=/etc/kubernetes/manifests/ \

保存

systemctl daemon-reload

systemctl restart kubelet

docker ps | grep webtool

22. Set configuration context \$ kubectl config use-context ik8s

In this task, you will configure a new Node, **ik8s-node-0**, to join a Kubernetes cluster as follows:

- Configure **kubelet** for automatic certificate rotation and ensure that both server and client CSRs are automatically approved and signed as appropriate via the use of **RBAC**.
- Ensure that the appropriate **cluster-info ConfigMap** is created and configured appropriately in the correct namespace so that future Nodes can easily join the cluster
- Your bootstrap **kubeconfig** should be created on the new Node at **/etc/kubernetes/bootstrap-kubelet.conf** (do not remove this file once your Node has successfully joined the cluster)
- The appropriate cluster-wide CA certificate is located on the Node at **/etc/kubernetes/pki/ca.crt**. You should ensure that any automatically issued certificates are installed to the node at **/var/lib/kubelet/pki** and that the **kubeconfig** file for **kubelet** will be rendered at **/etc/kubernetes/kubelet.conf** upon successful bootstrapping
- Use an additional group for bootstrapping Nodes attempting to join the cluster which should be called **system:bootstrappers:cka:default-node-token**
- Solution should start automatically on boot, with the **systemd** service unit file for **kubelet** available at **/etc/systemd/system/kubelet.service**

To test your solution, create the appropriate resources from the spec file located at **/opt/....kubernetes/kube-flannel.yaml**. This will create the necessary supporting resources as well as the **kube-flannel-ds DaemonSet**. You should ensure that this **DaemonSet** is correctly deployed to the single node in the cluster.

Hints:

- **kubelet** is not configured or running on **ik8s-master-0** for this task, and you should not

attempt to configure it.

- You will make use of TLS bootstrapping to complete this task.
- You can obtain the IP address of the Kubernetes API server via the following command `$ ssh ik8s-node-0 getent hosts ik8s-master-0`
- The API server is listening on the usual port, 6443/tcp, and will only server TLS requests
- The `kubelet` binary is already installed on `ik8s-node-0` at `/usr/bin/kubelet` . You will not need to deploy `kube-proxy` to the cluster during this task.
- You can `ssh` to the new worker node using `$ ssh ik8s-node-0`
- You can `ssh` to the master node with the following command `$ ssh ik8s-master-0`
- No further configuration of control plane services running on `ik8s-master-0` is required
- You can assume elevated privileges on both nodes with the following command `$ sudo -i`
- Docker is already installed and running on `ik8s-node-0`

Question weight: 8%

23. Set configuration context `$ kubectl config use-context bk8s`

Given a partially-functioning Kubernetes cluster, identify symptoms of failure on the cluster.

Determine the node, the failing service and take actions to bring up the failed service and restore the health of the cluster. Ensure that any changes are made permanently.

The worker node in this cluster is labelled with `name=bk8s-node-0`

Hints:

- You can `ssh` to the relevant nodes using `$ ssh $(NODE)` where `$(NODE)` is one of `bk8s-master-0` or `bk8s-node-0`
- You can assume elevated privileges on any node in the cluster with the following command `$ sudo -i`

Question weight: 4%

是 `kube-manager-controller` 没有启动 启动就做完了

`kc get cs`

能看到 `controller manager` 没有启动

登陆到 `master` 上

`systemctl start kube-manager-controller.service` # 到机器上的
`/etc/systemd/system` 确认一下是不是这么写的

24. Set configuration context `$ kubectl config use-context hk8s`

Creae a persistent volume with name `app-config` of capacity `1Gi` and access mode `ReadWriteOnce`. The type of volume is `hostPath` and its location is `/srv/app-config`

Question weight: 3%

Answer:

`apiVersion: v1`

`kind: PersistentVolume`

metadata:

name: app-config

spec:

capacity:

storage: 1Gi

accessModes:

- ReadWriteOnce

hostPath:

path: /srv/app-config